

SMART CURRENCY DETECTION AND TRANSACTION VERIFICATION FOR VISUALLY IMPAIRED

PHASE I REPORT

Submitted by

DHAANYA GK(2116230701069)

FAREED AHAMED KM (2116230701087)

S SANJAY KRISHNA (2116230701288)

SAUNAK RAMIYA SEBASAN (2116230701298)

*in partial fulfillment for the award of the degree
of*

BACHELOR OF ENGINEERING

IN

COMPUTER SCIENCE AND ENGINEERING



RAJALAKSHMI ENGINEERING COLLEGE, CHENNAI

ANNA UNIVERSITY: CHENNAI 600 025

OCTOBER 2026

ANNA UNIVERSITY: CHENNAI 600 025

BONAFIDE CERTIFICATE

Certified that this project report “**SMART CURRENCY DETECTION AND TRANSACTION VERIFICATION FOR VISUALLY IMPAIRED**” is the bonafide work of “**DHAANYA GK (2116230701069), FAREED AHAMED KM (2116230701087), S SANJAY KRISHNA (2116230701288) and SAUNAK RAMIYA SEBASAN (2116230701298)**” who carried out the project work under my supervision.

This Project addresses the following Sustainable Development Goal(s):

Goal-10: Reduced Inequality

SIGNATURE

Dr. J. Manoranjini, M.E., Ph.D.,
Professor and HOD,
Department of Computer Science and
Engineering,
Rajalakshmi Engineering College,
Chennai – 602105

SIGNATURE

Dr. Chettiyar Vani. V, M.E., Ph.D.
Professor
Department of Computer Science and
Engineering
Rajalakshmi Engineering College,
Chennai - 602105

Submitted for the Project Viva-Voce examination held on _____

INTERNAL EXAMINER

EXTERNAL EXAMINER

ABSTRACT

Cash remains the dominant mode of transaction in cash-driven economies such as India, yet it offers visually impaired users no independent way to confirm the denomination of a currency note. This forces reliance on a sighted companion or on the honesty of the other party, a recurring barrier to financial independence. This project develops an offline assistive system that combines computer vision and text-to-speech to close this gap. A captured note image is passed to a smartphone application, which applies OpenCV preprocessing followed by a YOLOv8-nano object detector trained on eleven classes covering both old and new series Indian currency notes. A confidence-threshold gate decides whether the prediction is reliable enough to announce or whether the note should be captured again; accepted detections are converted to speech by an on-device Piper TTS engine and delivered to the user over Bluetooth audio, with no dependency on internet connectivity at any stage. The detector was trained for 100 epochs on a CPU-only workstation and achieved a precision of 0.980, a recall of 1.000 and a mAP@0.5 of 0.995 on the validation split, with an inference latency of approximately 134 milliseconds per frame.

ACKNOWLEDGEMENT

Initially we thank the Almighty for being with us through every walk of our life and showering his blessings through the endeavor to put forth this report. Our sincere thanks to our Chairman **Mr. S. MEGANATHAN, B.E, F.I.E.,** our Vice Chairman **Mr. ABHAY SHANKAR MEGANATHAN, B.E., M.S.,** and our respected Chairperson **Dr. (Mrs.) THANGAM MEGANATHAN, Ph.D.,** for providing us with the requisite infrastructure and sincere endeavoring in educating us in their premier institution.

Our sincere thanks to **Dr. S.N. MURUGESAN, M.E., Ph.D.,** our beloved Principal for his kind support and facilities provided to complete our work in time. We express our sincere thanks to **Dr. J. MANORANJINI, M.E., Ph.D.,** Professor and Head of the Department of Computer Science and Engineering for her guidance and encouragement throughout the project work. We convey our sincere and deepest gratitude to our internal guide, **Dr. CHETTIYAR VANI VIVEKANAND, M.E., Ph.D.,** Professor, Department of Computer Science and Engineering for her valuable guidance and useful tips during our review throughout the course of the project.

SAUNAK RAMIYA SEBASAN: 230701298

S SANJAY KRISHNA: 230701288

FAREED AHAMED KM: 230701087

DHAANYA GK: 230701069

TABLE OF CONTENTS

	ABSTRACT	
	LIST OF TABLES	
	LIST OF FIGURES	
	LIST OF ABBREVIATIONS	
1.	INTRODUCTION	1
1.1	General	1
1.2	Motivation	1
1.3	Objectives	2
1.4	Scope of the Project	2
1.5	Organization of the Report	3
2.	LITERATURE SURVEY	4
2.1	General	4
2.2	Summary of Reviewed Work	4
2.3	Research Gap	6
3.	SYSTEM ANALYSIS	8
3.1	Existing System	8
3.2	Limitations of the Existing System	8
3.3	Proposed System	8
3.4	Problem Statement	9
4.	SYSTEM REQUIREMENTS	10
4.1	Hardware Requirements	10
4.2	Software Requirements	10
5.	SYSTEM DESIGN	12
5.1	System Architecture	12
5.2	Process Flow	12
5.3	Module Description	13
5.4	Project Structure	13
6.	METHODOLOGY	15
6.1	Dataset Preparation	15

6.2	Model Selection	15
6.3	Training Configuration	15
6.4	Confidence-Threshold Recapture Logic	16
6.5	Offline Voice Feedback	16
6.6	Offline Operation	17
7.	IMPLEMENTATION AND RESULT DISCUSSION	18
7.1	Training Loss Trend	18
7.2	Precision and Recall Progression	19
7.3	mAP Progression	19
7.4	Per-Class Detection Performance	20
7.5	Overfitting Assessment	23
7.6	Inference Latency	24
7.7	Comparison with Literature	24
7.8	Current Development Status	24
8.	CONCLUSION AND FUTURE WORK	26
8.1	Conclusion	26
8.2	Future Work	26
	REFERENCES	28

LIST OF TABLES

2.1	Summary of Literature Survey	4
4.1	Hardware Requirements and Approximate Cost	10
4.2	Software / Technology Stack	10
5.1	Module Description	13
6.1	YOLOv8n Training Configuration	15
7.1	Training Loss Trend (Epoch 1 / 50 / 100)	18
7.2	Precision Progression	19
7.3	Recall Progression	19
7.4	mAP@0.5 Progression	19
7.5	mAP@0.5:0.95 Progression	20
7.6	Per-Class Detection Performance	20
7.7	Inference Latency Breakdown	24
7.8	Development Status Checklist	25

LIST OF FIGURES

5.1	System Architecture – Wearable Indian Currency Recognition System	12
7.1	Training and Validation Loss and Metric Curves over 100 Epochs	18
7.2	Confusion Matrix on the Validation Set	21
7.3	Ground-Truth Labels – Validation Batch Sample	22
7.4	Model Predictions – Validation Batch Sample	23

LIST OF SYMBOLS, ABBREVIATIONS AND NOMENCLATURE

YOLO	– You Only Look Once (real-time object detection architecture)
TTS	– Text-to-Speech
mAP	– mean Average Precision
IoU	– Intersection over Union
CNN	– Convolutional Neural Network
FPS	– Frames Per Second
ONNX	– Open Neural Network Exchange
TFLite	– TensorFlow Lite
CPU	– Central Processing Unit
GPU	– Graphics Processing Unit
NPU	– Neural Processing Unit
SDG	– Sustainable Development Goal
dfl	– Distribution Focal Loss

CHAPTER 1

INTRODUCTION

1.1 General

Assistive technology for the visually impaired has traditionally focused on navigation and obstacle avoidance, with comparatively little attention paid to a task that sighted individuals perform without thought several times a day: handling cash. Cash remains a dominant mode of transaction worldwide, and in cash-driven economies such as India it continues to be the default choice for a very large share of everyday purchases. For a visually impaired person, a banknote carries no information that can be perceived by touch alone – Indian currency notes of different denominations are similar in size and texture – and there is presently no independent, on-the-spot way to confirm a note's denomination.

This project, undertaken by Team T032 of the Department of Computer Science and Engineering, Rajalakshmi Engineering College, proposes a camera-based assistive system that recognizes Indian currency denominations and communicates the result to the user through speech. The captured currency frame is processed on a companion smartphone, and the detection and offline text-to-speech pipeline works without any dependency on an internet connection.

1.2 Motivation

In cash-driven economies such as India, the inability to independently verify currency is a recurring, error-prone barrier to financial independence. It forces a visually impaired user to rely on a sighted companion or on the honesty of a shopkeeper or counterpart in the transaction, and the exposure to error is greatest at precisely the point where the mistake would be hardest for the user to catch. Existing assistive currency-recognition tools, reviewed in Chapter 2,

largely identify a note and speak its denomination. This project builds a low-cost, fully offline version of that core capability, using a lightweight detector that runs on a smartphone and a confidence check that avoids announcing uncertain results.

1.3 Objectives

- To design and train a lightweight, real-time object-detection model capable of recognizing Indian currency notes across both old and new note series.
- To preprocess a captured currency frame on the smartphone and run the trained detector on-device, without dependence on any cloud service.
- To deliver the detected denomination to the user as speech through a fully offline text-to-speech engine and Bluetooth audio output, so that the system has no dependency on internet connectivity.
- To incorporate a confidence-threshold recapture mechanism so that uncertain detections are re-checked rather than reported.

1.4 Scope of the Project

This report covers the first phase of the project, which delivers the core recognition capability of the system.

The scope of this phase covers dataset preparation, training and evaluation of a YOLOv8-nano detector, confidence-based acceptance of predictions, and voice announcement of the detected denomination through an offline TTS engine and Bluetooth audio.

Counterfeit or authenticity detection is explicitly excluded from the scope of this project, as is the use of UV/IR imaging hardware and expensive commercial smart-glasses hardware. The system logic is deliberately kept

simple and rule-based – a confidence threshold decides whether a prediction is announced – based on the reasoning that additional complexity increases cognitive load for assistive-technology users.

1.5 Organization of the Report

The remainder of this report is organized as follows. Chapter 2 reviews the existing literature on assistive currency-recognition systems. Chapter 3 analyses the existing systems and states the problem being addressed. Chapter 4 lists the hardware and software requirements. Chapter 5 presents the system architecture, process flow and module description. Chapter 6 describes the methodology followed for dataset preparation, model training and voice feedback. Chapter 7 presents and discusses the implementation results, including the YOLOv8n training outcomes. Chapter 8 concludes the report and outlines directions for further improvement.

CHAPTER 2

LITERATURE SURVEY

2.1 General

This chapter surveys twenty recent works on assistive currency recognition and related vision-based accessibility systems [1]–[20], listed in Table 2.1 in alphabetical order of the first author’s surname. The survey covers detector architectures (CNN [3], [5], [9], [14], YOLOv5 [11], [15], YOLOv8 [7], YOLOv11 [4], YOLO-NAS [2], EfficientDet [1], MobileNet/EfficientNet [18]), feedback mechanisms (TTS and audio [7], [14], [15], [16], [17], tactile [3]), and adjacent problems such as counterfeit detection [2], [6], [10], [12], [13], [19] and multimodal assistive agents [8], [20], in order to position the objectives of this project against the current state of the art.

2.2 Summary of Reviewed Work

Table 2.1 summarises the methodology and key inference of each reviewed work.

Table 2.1 Summary of Literature Survey

No.	Author(s) / Year	Title	Methodology	Inference
1	Bandara (2025) [1]	Money Recognition for the Visually Impaired: A Case Study on Sri Lankan Banknotes	EfficientDet object detection with image preprocessing.	Accurate banknote recognition; establishes a regional benchmark.
2	Chhabra, Goyal (2025) [2]	A Two-Phase Deep Learning Model for Counterfeit Detection of Indian Banknotes using YOLO-NAS and UV Imaging	Two-phase pipeline using YOLO-NAS and UV imaging for security-feature localization.	Detects denomination and counterfeit currency using security-ink features not visible under normal lighting.
3	Edupuganti, Sneha, Vardhan, Kumar (2024) [3]	Generative AI-Based Currency Detector for Visually Impaired	CNN-based image classification with high-resolution camera, tactile buttons, voice commands and audio output.	Recognizes denomination and announces it through audio; tactile controls improve accessibility.
4	Fanzury, Hwang (2026) [4]	Real-Time Identification of Mixed and Partly Covered	YOLOv11 detection with hybrid TFLite / server	Precision 0.951, recall 0.935, mAP@0.5

No.	Author(s) / Year	Title	Methodology	Inference
		Foreign Currency Using YOLOv11 Object Detection	inference; multi-currency and occlusion handling.	0.965; handles mixed and partly covered currencies effectively.
5	Ghanem, Youness, Wahba, Abdelaal (2025) [5]	Recognizing Egyptian Currency for People with Visual Impairment Using Deep Learning Models	CNN-based classification on a regional Egyptian currency dataset.	Provides denomination recognition and addresses a regional gap in accessibility tools.
6	Hantoush (2025) [6]	Artificial Intelligence Model of Iraqi Paper Currency Detection in Vending Machines	CNN classification with embedded deployment and microscopic-level feature analysis.	Combined denomination and authenticity verdict for vending/POS integration.
7	Lalitha, Gayatri, Selvi, Saranya, Ponugoti (2025) [7]	Assisting Visually Challenged Individuals to Distinguish Authentic Currency from Counterfeit	YOLOv8 detection with GTTS speech synthesis; Roboflow annotation and standard detection-metric evaluation.	Provides spoken denomination, currency type and validity in a single audio-based pipeline.
8	Maheswari, Aluvalu, Dhanaraj et al. (2026) [8]	An Assistive Framework for Visually Impaired Using Deep Learning and Natural Language Processing	CNN/RNN/LSTM pipeline with NLP caption generation and speech synthesis.	Narrative spoken scene descriptions including currency identification.
9	Nasayreh, Jaradat, Gharaibeh et al. (2024) [9]	Jordanian Banknote Data Recognition: A CNN-Based Approach with Attention Mechanism	CNN with an attention mechanism and attention-guided feature extraction.	High-accuracy denomination recognition; attention improves feature extraction and robustness.
10	Ovi, Hossain, Biswas (2025) [10]	Quantitative Currency Evaluation in Low-Resource Settings through Pattern Analysis	Lightweight custom CNN with feature-based template matching and a Unified Currency Damage Index.	Identifies denomination, provides a usability/damage score and a counterfeit flag in a compact on-device pipeline.
11	Pasumarthy, Padhy, Yadav, Subramaniam, Rao (2022) [11]	An Indian Currency Recognition Model for Assisting Visually Impaired Individuals	YOLOv5 detection on an annotated, augmented dataset of ~10,000 images (224×224), trained under cluttered, reflective and blurred conditions.	92.71% accuracy on unseen data; robust to reflections, blur and clutter; supported Raspberry Pi and smartphone deployment.
12	Pham, Park, Nguyen, Batchuluun, Park (2020) [12]	Deep Learning-Based Fake-Banknote Detection for the Visually Impaired Using Smartphone Cameras	CNN-based counterfeit detection using RGB image analysis.	Identifies counterfeit banknotes using ordinary smartphone cameras.

No.	Author(s) / Year	Title	Methodology	Inference
13	Pham, Lee, Park, Park (2022) [13]	Deep Learning-Based Fake Multinational Banknotes Detection in a Cross-Dataset Environment	YOLOv3 detection combined with ResNet-18, GoogleNet and Inception-ResNet-v2 fusion.	99.289% mAP for counterfeit classification; strong cross-currency generalization.
14	Puthumalar, Sivaranjani (2024) [14]	Currency Recognition for Visually Impairment People Using Deep Learning Algorithm	CNN with image preprocessing and audio feedback.	Lightweight, real-time denomination recognition designed for accessibility.
15	Sathwik, Balaji, Chandralekha (2026) [15]	Real-Time Currency Identification for Visually Impaired Users Using a Lightweight YOLOv5-Based Object Detection System	YOLOv5-Nano detection with live camera-feed processing and audio feedback.	99.5% mAP at IoU 0.5; low-latency edge inference suitable for real-time assistive applications.
16	Shreya, Siddique, Tasnim (2025) [16]	Real-Time Currency Detection and Voice Feedback for Visually Impaired Individuals	Real-time object detection combined with live voice feedback.	Spoken denomination feedback emphasising real-time responsiveness for daily use.
17	Siddappa, Gowda, Hasham, Kumar (2026) [17]	AI-Based Money Recognition Assistant for Blind People	Lightweight CNN on TensorFlow Lite deployed on an ESP32-CAM with INT8 quantization.	93.8% accuracy, ~1.2 s inference; combines currency recognition with obstacle assistance in an offline wearable system.
18	Subreena, Hossain, Raquib, Murad, Prity, Hanif, Rahimi (2026) [18]	Robust and Real-Time Bangladeshi Currency Recognition: A Dual-Stream MobileNet and EfficientNet Approach	Dual-stream MobileNet and EfficientNet architecture with Explainable AI.	Real-time recognition using an efficient hybrid architecture suitable for edge devices.
19	Wasnik, Temurnikar (2026) [19]	Deep Learning-Based Spectral-Spatial Model for Counterfeit Currency Detection	Dual-branch CNN with attention-guided fusion of RGB and UV/IR imaging.	Robust counterfeit classification under low light, blur and noise.
20	Zhao, Wang, Geng, Yu, Li (2025) [20]	“Less is More”: Reducing Cognitive Load and Task Drift in Real-Time Multimodal Assistive Agents	Multimodal agent design and evaluation with real-time feedback and task-drift analysis.	Design guidelines for reducing cognitive load and task drift in real-time assistive agents.

2.3 Research Gap

Across the reviewed literature, YOLO-family detectors (YOLOv5/v8/v11) consistently achieve high denomination-recognition accuracy [4], [5], [7], [11], [15] – several works report mAP@0.5 above 0.95 [4], [5], [15] – and are

increasingly deployed on lightweight, edge-friendly formats such as TFLite [4], [17], making them well suited to smartphone or microcontroller-adjacent deployment. Text-to-speech and audio feedback are near-universal in this space [3], [7], [14], [15], [16], [17], confirming that voice output is the established modality for communicating detection results to visually impaired users.

However, every system surveyed treats each currency note as an isolated recognition event: the note is identified, its denomination is spoken, and the interaction ends. Additionally, several counterfeit-detection works (rows 2, 12, 13 and 19) [2], [12], [13], [19] rely on UV/IR imaging hardware or cross-dataset fusion models that add cost and computational overhead disproportionate to the everyday goal of denomination recognition; this project deliberately keeps counterfeit detection out of scope in favour of a low-cost, fully offline, rule-based design.

CHAPTER 3

SYSTEM ANALYSIS

3.1 Existing System

Current assistive currency-recognition approaches, as surveyed in Chapter 2, fall into three broad categories: (i) CNN or YOLO-based denomination detectors deployed on smartphones or single-board computers such as Raspberry Pi [1], [4], [5], [9], [11], [14], [15], [18]; (ii) GTTS or other cloud-dependent TTS pipelines paired with a detection model [7]; and (iii) spectacle-mounted or wearable prototypes that combine currency recognition with general obstacle detection [17]. These systems are generally able to identify the note currently in view and speak its denomination with reasonably high accuracy [11], [14], [16].

3.2 Limitations of the Existing System

- Several systems rely on cloud-based TTS (e.g., gTTS) [7], which requires an active internet connection and is unsuitable for the fully offline operating requirement of this project.
- Counterfeit-detection variants require additional UV/IR hardware [2], [19], increasing cost and system complexity beyond what is needed for the core denomination-verification use case.
- Most prototypes are demonstrated as standalone recognition tools [6], [10], [16], some with online or computationally heavy components [8], [13], rather than as a lightweight, fully offline pipeline suited to on-phone inference.

3.3 Proposed System

The proposed system addresses these gaps with a lightweight, fully offline design. A currency frame captured by the camera is passed to a smartphone application, which performs OpenCV preprocessing and YOLOv8-nano

inference, applies a confidence threshold to decide whether a detection is reliable, and – if it is – converts the denomination to speech using an offline Piper TTS engine, delivered to the user over Bluetooth audio. If confidence is below the threshold, the system does not announce the result and asks for the note to be captured again instead of reporting an uncertain result.

3.4 Problem Statement

Visually impaired users have no independent way to confirm a currency note's denomination. In cash-driven economies such as India, this is a recurring, error-prone barrier at exactly the point where mistakes are hardest for a visually impaired user to catch. This project aims to close the recognition part of that gap with a low-cost, fully offline system that performs real-time denomination detection and announces the result by voice.

CHAPTER 4

SYSTEM REQUIREMENTS

4.1 Hardware Requirements

The prototype is intentionally designed around hardware that is already commonly available: a workstation for training the model, a smartphone for on-device inference, and Bluetooth earphones or a speaker for audio output. Table 4.1 lists the hardware components.

Table 4.1 Hardware Requirements and Approximate Cost

Component	Use	Approx. Cost
Training workstation (12th Gen Intel Core i7-1255U, CPU only)	Dataset preparation, model training and evaluation	Already available
Smartphone	Local inference (OpenCV, YOLO, TTS)	Already owned
Bluetooth earphones / speaker	Audio feedback to the user	Already owned

4.2 Software Requirements

Table 4.2 lists the software and technology stack used in this work. Python environment and dependency management is handled exclusively through uv, with pyproject.toml as the single source of truth and uv.lock committed for reproducibility. Docker is used strictly to standardize the training/evaluation environment across the team and is not used on the phone-side inference app, which cannot run containers.

Table 4.2 Software / Technology Stack

Stage	Tool / Technology	Purpose
Image input	Captured currency frame (image)	Input to the smartphone application
Preprocessing	Python 3.11 + OpenCV	Resize, crop, normalize, denoise, brightness enhancement
Detection	YOLOv8-nano (Ultralytics)	Currency denomination detection
Training	Python 3.11 + Ultralytics / PyTorch	Model training and validation

Stage	Tool / Technology	Purpose
Environment / dependency mgmt.	uv (pyproject.toml + uv.lock)	Reproducible Python environment
Dev / training containerization	Docker + docker-compose	Standardized training / evaluation environment only
Confidence logic	Python rule-based threshold logic	Accept reliable predictions / request recapture
Voice output	Piper TTS (offline)	Convert detected denomination to speech
Audio delivery	Bluetooth earphones / speaker	Deliver spoken result to the user
Code quality	ruff format, ruff check, type hints	Formatting, linting and static typing

CHAPTER 5

SYSTEM DESIGN

5.1 System Architecture

Figure 5.1 shows the overall architecture of the system. A captured currency frame is passed to the smartphone application, whose processing system contains two blocks: Currency Detection (image preprocessing, YOLO detection of the denomination, and confidence scoring using the quantized model) and Voice Feedback (label mapping of the denomination to text, on-device text-to-speech, and generation of a user-friendly voice message). The generated audio is delivered to the output unit, where Bluetooth earphones or a speaker announce the denomination, for example “Rs. 500”, to the user.




Figure 5.1 System Architecture – Wearable Indian Currency Recognition System

5.2 Process Flow

The process flow of the system is as follows:

1. Image input – a camera captures the currency note and the frame is provided to the smartphone application.
2. Preprocessing – OpenCV resizes and normalizes the image and applies basic brightness enhancement or denoising.

3. Denomination detection – the YOLOv8n model predicts the currency class, bounding box and confidence score.
4. Confidence handling – a confidence threshold determines whether the prediction is reliable; a low-confidence result triggers a request to capture the note again.
5. Voice output – the denomination label is mapped to text and the Piper TTS engine converts it into speech.
6. User feedback – the result is delivered through Bluetooth earphones or a speaker.

5.3 Module Description

The system is organized into two modules, summarised in Table 5.1.

Table 5.1 Module Description

Module	Function	Key Components
M1 – Currency Detection	YOLO-based denomination detection from a captured frame	OpenCV preprocessing, trained/quantized YOLO model, confidence scoring
M2 – Voice Feedback	TTS output announcing the detected denomination	Label-to-speech mapping, on-device Piper TTS engine, Bluetooth audio output

Inference runs entirely on the connected mobile phone, and no component of the pipeline requires an internet connection.

5.4 Project Structure

The codebase follows a modular structure in which the detection and audio modules are independently runnable and demoable. Scripts only orchestrate training, evaluation and deployment entry points, while all core logic resides in dedicated source modules for detection, audio, configuration and shared utilities. Model hyperparameters and thresholds are centralized in a

configuration module, and the Python environment is managed with uv for reproducibility.

CHAPTER 6

METHODOLOGY

6.1 Dataset Preparation

The detector was trained on an annotated Indian currency note dataset covering eleven classes spanning both the old and new note series (e.g., ₹10, ₹20, ₹50, ₹100, ₹200, ₹500 and ₹2000, each tagged by series where applicable). The dataset was split into 459 training images and 23 validation images, formatted for YOLO training with a data.yaml configuration file. Class mappings and split paths were verified as the first step before training began.

6.2 Model Selection

YOLOv8-nano was selected as the detection backbone specifically for its low latency on phone-side inference: at roughly 3 million parameters and 8.2 GFLOPs, it is small enough to run efficiently on a mid-range smartphone CPU/NPU while still supporting transfer learning from COCO-pretrained weights. All model hyperparameters – image size, confidence threshold, IoU threshold and the class list – are centralized in a dedicated configuration module rather than hardcoded inline, so that they can be re-tuned from the actual training results without touching the detection or audio code.

6.3 Training Configuration

The model was trained for 100 epochs on a CPU-only workstation (12th Gen Intel Core i7-1255U, no GPU), using the Ultralytics YOLOv8 training pipeline with an image size of 640×640 and a batch size of 16. Table 6.1 summarises the key training configuration parameters.

Table 6.1 YOLOv8n Training Configuration

Parameter	Value
Base model	YOLOv8n (nano)
Framework	Ultralytics 8.4.126, PyTorch 2.13.0 (CPU build)
Hardware used for training	12th Gen Intel Core i7-1255U (CPU only, no GPU)
Python version	3.11.13
Number of classes	11 (old and new series Indian currency notes)
Training images	459
Validation images	23
Image size	640 × 640
Batch size	16
Epochs	100
Optimizer	AdamW (auto-selected), lr0 = 0.000667, momentum = 0.9
Early-stopping patience	20 epochs (never triggered)
Augmentation	HSV jitter, horizontal flip (p = 0.5), mosaic, random erasing
Total training time	5.147 hours
Model size	130 layers, 3,012,993 parameters, 8.2 GFLOPs

6.4 Confidence-Threshold Recapture Logic

Rather than reporting every detection, the smartphone application applies a confidence threshold (tuned from the training results and stored in the configuration module) to decide whether a prediction should be spoken. If the top detection's confidence is below the threshold, the system asks the user to capture the note again instead of announcing an uncertain denomination – this avoids the failure mode of confidently mis-stating a note's value to a user who cannot visually double-check the result.

6.5 Offline Voice Feedback

Text-to-speech is implemented using Piper TTS, chosen because it is fully offline, lightweight and provides good voice quality through a simple Python API, in line with the project's strict offline-only constraint. Cloud or API-

based TTS engines such as gTTS were explicitly avoided because they require an active internet connection, which the system is designed never to depend on.

6.6 Offline Operation

The complete pipeline – preprocessing, inference, confidence handling and speech synthesis – runs on the smartphone without any network connection. No tool, API or library that assumes internet access is used anywhere in the pipeline, including the text-to-speech engine, so the system behaves identically with or without connectivity.

CHAPTER 7

IMPLEMENTATION AND RESULT DISCUSSION

The YOLOv8n detector was trained for the full 100 epochs on a CPU-only workstation; the results below are drawn directly from the training run's console log and per-epoch validation metrics. Figure 7.1 plots the training/validation loss components and the precision, recall and mAP metrics across all 100 epochs, discussed section by section below.




Figure 7.1 Training and Validation Loss and Metric Curves over 100 Epochs

7.1 Training Loss Trend

Table 7.1 Training Loss Trend (Epoch 1 / 50 / 100)

Loss	Epoch 1	Epoch 50	Epoch 100	Trend
box_loss	1.252	0.531	0.266	Steady decline – model learns to localize bounding boxes
cls_loss	3.914	0.520	0.302	Steep early drop, gradual convergence – class separation improves
dfl_loss	1.724	1.112	0.991	Slow, steady – distribution focal loss refines box edges

box_loss fell 79% across training, reflecting steadily improving bounding-box localization; cls_loss fell 92%, the dominant early loss as the model initially

struggled to separate 11 visually similar currency classes before learning quickly; dfl_loss declined more gradually (43%), consistent with its role in refining fine-grained box boundaries.

7.2 Precision and Recall Progression

Table 7.2 Precision Progression

Phase	Value	Interpretation
Epoch 1	0.008	Near-zero – model predicts everything, most predictions wrong
Epoch 2	0.918	Sharp jump – model learns to be selective
Epoch 3–10	0.22 – 0.91	Volatile – model oscillates between cautious and aggressive
Epoch 11+	0.88 – 0.98	Stable and high – model is confident and correct
Best (epoch 87)	0.980	98% of detected objects are actual currency notes

Table 7.3 Recall Progression

Phase	Value	Interpretation
Epoch 1	1.000	Detects everything (but mostly wrong)
Epoch 2	0.050	Collapses – model becomes overly conservative
Epoch 3–10	0.40 – 0.86	Recovers steadily
Epoch 11+	0.93 – 1.00	Catches nearly all currency notes
Best (epoch 87)	1.000	Finds every currency note in the validation set

The epoch-2 recall collapse coincides with the sharp precision jump the same epoch – the model swung from “predict everything” to “predict almost nothing” before settling into a stable balance from epoch 11 onward.

7.3 mAP Progression

Table 7.4 mAP@0.5 Progression

Phase	Value	Interpretation
Epoch 1	0.266	Low – model barely overlaps ground truth
Epoch 3	0.610	Rapid improvement
Epoch 7	0.878	Crosses into production-grade territory
Epoch 11	0.995	Ceiling reached

Phase	Value	Interpretation
Epochs 12–100	0.962 – 0.995	Stays at or near ceiling for the remaining 89 epochs

Table 7.5 mAP@0.5:0.95 Progression

Phase	Value	Interpretation
Epoch 1	0.183	Very loose detections
Epoch 11	0.727	Good, but boxes not pixel-precise
Epoch 44	0.883	Boxes getting tighter
Epoch 75	0.905	Near-optimal box precision
Best (epoch 87)	0.910	91% overlap accuracy at strict IoU thresholds

mAP@0.5 reached its 0.995 ceiling by epoch 11 and stayed there for the remaining 89 epochs, while the stricter mAP@0.5:0.95 metric kept improving all the way to epoch 87 – evidence that the model continued refining bounding-box precision on objects it had already learned to detect, rather than merely memorizing the training set.

7.4 Per-Class Detection Performance

All eleven classes reached mAP@0.5 = 0.995; the variance across classes shows up only in the stricter mAP@0.5:0.95 metric, summarised in Table 7.6.

Table 7.6 Per-Class Detection Performance (mAP@0.5:0.95)

Class	mAP@0.5:0.95	Notes
₹50 (New)	0.995	Best – single instance, model nailed it
₹10 (Old)	0.951	Strong – distinct visual features
₹20 (Old)	0.945	Strong
₹2000	0.895	Average
₹200	0.895	Average
₹20 (New)	0.895	Average
₹500 (New)	0.895	Average
₹50 (Old)	0.895	Average
₹100 (New)	0.889	Slightly below average
₹100 (Old)	0.846	Lowest – only 2 validation images, likely harder variations

No class performs significantly worse than the others; the spread (0.846–0.995) is driven mainly by the small number of validation images available per class (1–4 images), rather than by any structural weakness in the model.

Figure 7.2 presents the confusion matrix computed on the validation split. Every predicted class lies on the diagonal, matching its true class with no off-diagonal confusions anywhere in the matrix, corroborating the precision and recall values reported above and confirming that the detector does not confuse visually similar denominations, such as the old- and new-series notes of the same face value.




Figure 7.2 Confusion Matrix on the Validation Set

Figures 7.3 and 7.4 show a representative validation batch with the ground-truth annotations and the corresponding model predictions overlaid on the same images. The predicted class labels and confidence scores closely match the ground-truth labels across all sixteen sample notes in the batch, including notes captured at an angle, partially folded, or photographed against varied backgrounds and lighting.



Figure 7.3 Ground-Truth Labels – Validation Batch Sample



Figure 7.4 Model Predictions – Validation Batch Sample

7.5 Overfitting Assessment

No significant overfitting was observed. Validation losses for all three loss components decreased steadily throughout training with no point at which validation loss rose while training loss continued to fall – the classic overfitting signature. $mAP@0.5$ plateaued at its ceiling from epoch 11 without ever dropping sharply and recovering, which would suggest memorization, and $mAP@0.5:0.95$ continued climbing for 76 further epochs, indicating genuine, ongoing refinement of box precision. YOLOv8n's small capacity (about 3 million parameters) relative to the 459-image training set further limits the model's ability to memorize the data even if it were inclined to. Early stopping (patience = 20) never triggered, since the model kept matching or beating its best $mAP@0.5$ well past epoch 80 – the ideal outcome of a model

that genuinely needed the full 100 epochs to converge rather than stalling early.

7.6 Inference Latency

Table 7.7 Inference Latency Breakdown (CPU, no GPU)

Stage	Latency
Preprocess	1.7 ms
Inference	131.0 ms
Postprocess	1.4 ms
Total (CPU)	~134 ms (~7.5 FPS)

An end-to-end latency of approximately 134 ms per frame (about 7.5 FPS) on a CPU-only workstation is acceptable for an assistive tool that announces a result once per capture rather than requiring continuous video-rate throughput. On a phone with GPU or NPU acceleration – the actual target deployment – this figure is expected to drop substantially, to roughly 20–40 ms per frame.

7.7 Comparison with Literature

The achieved validation performance (precision 0.980, recall 1.000, mAP@0.5 0.995) compares favourably with the strongest results reported in the literature survey of Chapter 2, several of which report mAP@0.5 in the 0.95–0.965 range using comparable YOLO-family detectors. This suggests that the YOLOv8n configuration and training regime used here are competitive with the current state of the art for Indian-currency denomination detection, while remaining a substantially smaller and faster model than the CNN-attention or dual-branch architectures used in some of the surveyed counterfeit-detection systems.

7.8 Current Development Status

Table 7.8 summarises progress against the development checklist at the time of this report.

Table 7.8 Development Status Checklist

No.	Item	Status
1	Dataset preparation	Done
2	Train YOLOv8n	Done
3	Export model (ONNX / TFLite)	Done
4	Detection module	Done
5	Config tuning from training results	Done
6	Audio module (Piper TTS)	Done
7	Detection–audio pipeline integration	Done
8	End-to-end test (captured frame → detect → speak)	Done

The dataset, trained detector, exported model, detection module, tuned configuration and audio module are complete. Integration of the detection and audio modules into a single pipeline, unit tests, the end-to-end test and documentation remain to be completed before this phase is considered finished.

CHAPTER 8

CONCLUSION AND FUTURE WORK

8.1 Conclusion

This report has presented the design, methodology and results of a fully offline assistive system for Indian currency denomination detection aimed at visually impaired users. A YOLOv8-nano detector trained on eleven currency classes achieved a validation precision of 0.980, a recall of 1.000 and a mAP@0.5 of 0.995 with no evidence of overfitting, and an end-to-end inference latency of approximately 134 ms on CPU-only hardware, competitive with the strongest results reported in the literature survey. The detection and offline text-to-speech modules are complete; pipeline integration, automated testing, the end-to-end demonstration and documentation remain as the immediate next steps.

8.2 Future Work

Future work will focus on strengthening the current system. The dataset will be enlarged with more images per class and a wider range of lighting, blur, occlusion and reflective-surface conditions, so that per-class performance can be evaluated on a larger validation set. Inference latency and power usage will be benchmarked on the actual target smartphone hardware, and the detection and voice modules will be packaged into a single mobile application. The robustness of the detector under low light, blur and partial occlusion will also be evaluated systematically.

Consistent with the project's design philosophy, the system will be kept deliberately simple and rule-based – a confidence threshold rather than a complex decision framework – based on the reasoning that added complexity increases cognitive load for assistive-technology users. Counterfeit or authenticity detection, UV/IR imaging hardware and expensive commercial

smart-glasses hardware will continue to be explicitly excluded from the project's scope.

REFERENCES

1. Akshaan Bandara, 2025, Money Recognition for the Visually Impaired: A Case Study on Sri Lankan Banknotes, arXiv preprint arXiv:2502.14267, February – 2025.
2. Payal Chhabra, Sonali Goyal, 2025, A Two-Phase Deep Learning Model for Counterfeit Detection of Indian Banknotes using YOLO-NAS and UV Imaging for Visually Impaired People, EVERGREEN – Joint Journal of Novel Carbon Resource Sciences & Green Asia Strategy, Volume 12, Issue 03, pp. 1633–1653, September – 2025.
3. Geetha Rani E, Guru Sneha B, Harsha Vardhan N, Sai Kumar B, Anusha D, Sandyarani Vadlamudi, 2024, Generative AI-Based Currency Detector for Visually Impaired, 2024 International Conference on Signal Processing and Advance Research in Computing (SPARC), Volume 1, pp. 1–6, September – 2024.
4. Nanda Fanzury, Mintae Hwang, 2025, Real-Time Identification of Mixed and Partly Covered Foreign Currency Using YOLOv11 Object Detection, AI, Volume 6, Issue 10, Article 241, September – 2025.
5. Ahmed M. Ghanem, Hassan A. Youness, Mohamed Wahba, Hammam M. Abdelaal, 2025, Recognizing Egyptian Currency for People with Visual Impairment Using Deep Learning Models, Scientific Reports, Volume 15, Article 34288, October – 2025.
6. Shymaa Akram Hantoush, 2026, Artificial Intelligence Model of Iraqi Paper Currency Detection in Vending Machines, Ibn AL-Haitham Journal For Pure and Applied Sciences, Volume 39, Issue 02, pp. 299–311, April – 2026.
7. R. Lalitha, S. Gayathri, S. Selvi, K. Saranya, Sreelekha Ponugoti, 2025, Assisting Visually Challenged Individuals to Distinguish Authentic Currency from Counterfeit, Advances in Data Science and Artificial Intelligence: Proceedings of ERCICA 2024, Volume 1, pp. 219–232, July – 2025.
8. V. Uma Maheswari, Rajanikanth Aluvalu, Rajesh Kumar Dhanaraj, Mahmoud Ahmad Al-Khasawneh, Hariprasath Manoharan, Shitharth Selvarajan, 2026, An Assistive Framework for Visually Impaired Using Deep Learning and Natural Language Processing, International Journal of Speech Technology, Volume 29, Article 58, July – 2026.
9. Ahmad Nasayreh, Ameera S. Jaradat, Hasan Gharaibeh, Waed Dawaghreh, Rabia Mehamad Al Mamlook, Yaqeen Alqudah, Qais Al-Na'amneh, Mohammad Sh. Daoud,

Hazem Migdady, Laith Abualigah, 2024, Jordanian Banknote Data Recognition: A CNN-Based Approach with Attention Mechanism, Journal of King Saud University – Computer and Information Sciences, Volume 36, Issue 04, Article 102038, April – 2024.

10. Md Sultanul Islam Ovi, Mainul Hossain, Md Badsha Biswas, 2025, Quantitative Currency Evaluation in Low-Resource Settings through Pattern Analysis to Assist Visually Impaired Users, arXiv preprint arXiv:2509.06331, September – 2025.
11. Madhav Pasumarthy, Rutvi Padhy, Raghuveer Yadav, Ganesh Subramaniam, Madhav Rao, 2022, An Indian Currency Recognition Model for Assisting Visually Impaired Individuals, 2022 IEEE International Conference on Recent Advances in Systems Science and Engineering (RASSE), pp. 123–127, November – 2022.
12. Tuyen Danh Pham, Chanhum Park, Dat Tien Nguyen, Ganbayar Batchuluun, Kang Ryoung Park, 2020, Deep Learning-Based Fake-Banknote Detection for the Visually Impaired People Using Visible-Light Images Captured by Smartphone Cameras, IEEE Access, Volume 8, pp. 63144–63161, 2020.
13. Tuyen Danh Pham, Young Won Lee, Chanhum Park, Kang Ryoung Park, 2022, Deep Learning-Based Detection of Fake Multinational Banknotes in a Cross-Dataset Environment Utilizing Smartphone Cameras for Assisting Visually Impaired Individuals, Mathematics, Volume 10, Issue 09, Article 1616, May – 2022.
14. T. Puthumalar, R. Sivaranjani, 2024, Currency Recognition for Visually Impairment People Using Deep Learning Algorithm, International Journal of Innovative Research in Computer and Communication Technology (IJIRCT), Volume 10, Issue 04, pp. 1–4, August – 2024.
15. Nethi Sathwik, Talasila Balaji, M. Chandralekha, 2026, Real-Time Currency Identification for Visually Impaired Users Using a Lightweight YOLOv5-Based Object Detection System, 2026 International Conference on Data Science, Agents and Artificial Intelligence (ICDSAAI 2026), pp. 1–6, March – 2026.
16. Saraf Anzum Shreya, MD. Abu Ismail Siddique, Sharaf Tasnim, 2025, Real-Time Currency Detection and Voice Feedback for Visually Impaired Individuals, arXiv preprint arXiv:2510.20267, October – 2025.
17. Dr. Manjunatha Siddappa, Madhan Gowda R S, Mohammed Hasham, Naveen Kumar N, 2026, AI-Based Money Recognition Assistant For Blind People, INTERNATIONAL

18. Subreena, Mohammad Amzad Hossain, Mirza Raquib, Saydul Akbar Murad, Farida Siddiqi Prity, Muhammad Hanif, Nick Rahimi, 2026, Robust and Real-Time Bangladeshi Currency Recognition: A Dual-Stream MobileNet and EfficientNet Approach, arXiv preprint arXiv:2602.07015, January – 2026.
19. Kumud Wasnik, Ankit Temurnikar, 2026, Deep Learning-Based Spectral-Spatial Model for Counterfeit Currency Detection, International Journal of Computer Applications, Volume 187, Issue 124, pp. 19–26, July – 2026.
20. Yi Zhao, Siqi Wang, Qiqun Geng, Erxin Yu, Jing Li, 2025, “Less is More”: Reducing Cognitive Load and Task Drift in Real-Time Multimodal Assistive Agents for the Visually Impaired, arXiv preprint arXiv:2511.00945, November – 2025.

An Edge-Deployable YOLOv8n-Based Indian Currency Recognition System with Offline Speech Assistance

Chettiyar Vani Vivekanand

*Computer Science and Engineering
Rajalakshmi Engineering College, Chennai,
India
chettiyarvanivivekanand@rajalakshmi.edu.in*

Dhaanya GK

*Computer Science and Engineering
Rajalakshmi Engineering College,
Chennai, India
230701069@rajalakshmi.edu.in*

Fareed Ahamed KM

*Computer Science and Engineering
Rajalakshmi Engineering College,
Chennai, India
230701087@rajalakshmi.edu.in*

S Sanjay Krishna

*Computer Science and Engineering
Rajalakshmi Engineering College, Chennai,
India
230701288@rajalakshmi.edu.in*

Saunak Ramiya Sebasan

*Computer Science and Engineering
Rajalakshmi Engineering College,
Chennai, India
230701298@rajalakshmi.edu.in*

Abstract—This paper presents a real-time currency denomination detection system based on the YOLOv8n object detector, which is designed to recognize not only different denominations but also old and new design variants of Indian currency notes, resulting in an eleven-class detection scheme. For visually impaired people, independently identifying the denomination of a currency note can be a persistent everyday challenge, especially when older and newer versions of the same denomination are in circulation because differences in size, color, and printed features can affect how reliably a recognition system can identify them. The model was trained on a custom-annotated dataset of 482 images and achieved a precision of 0.98, recall of 1.00, and mAP@0.5 of 0.995 on the validation set. On CPU-only hardware, the system has a total inference latency of about 134 milliseconds per image, indicating its potential for low-cost deployment without dedicated GPU acceleration. An offline text-to-speech engine is used to convert the detected denomination into spoken output, which is especially useful for assistive applications. Experimental testing on unseen images consistently produced high-confidence detections with an average of 0.96. These findings demonstrate that a lightweight model can be used to achieve fine-grained recognition of currency note variants while maintaining real-time performance. As a result, the suggested system offers a workable and repeatable basis for assistive currency-recognition applications that can be implemented on edge devices with limited resources.

Index Terms—assistive technology, currency recognition, YOLOv8, object detection, text-to-speech, visually impaired, edge inference

I. INTRODUCTION

A significant portion of the world's population still suffers from visual impairment, and one of the many daily tasks made more challenging by this condition is identifying paper money. Unlike

digital payments, physical banknotes do not offer a direct way for a visually impaired person to independently confirm their denomination without outside assistance. This problem is made more difficult in currencies like the Indian Rupee, where older and newer design variants of the same denomination may continue to circulate at the same time, differing in size, color scheme, and printed security features. Because they offer a good balance between detection speed and accuracy, YOLO-based object detectors have become particularly useful for this type of application.

Previous research has explored currency denomination recognition extensively using deep learning techniques; systems based on YOLOv5 have demonstrated strong performance for Indian currency recognition, including under cluttered and visually degraded conditions [1], and lighter YOLOv5 Nano variants have demonstrated that real-time inference can be achieved on constrained hardware, making them suitable for edge-based deployment [2]. YOLOv8 and speech synthesis have been combined in more recent work to combine currency detection and audio feedback into a single low-latency pipeline, enabling the communication of denomination and validity information through spoken output [3]. This YOLOv8-plus-speech approach is especially conceptually similar to the system described in this paper, and thus serves as a main point of comparison.

Despite these advancements, the majority of current methods define currency classes based only on face value and do not specifically differentiate between different design generations of the same denomination. This distinction is crucial for currencies that are periodically redesigned because older and newer notes may

have distinctly different visual characteristics while still representing the same denomination; as a result, a model trained solely on current-series notes may perform unreliably when presented with older notes that are still in circulation, and the same problem may arise in the opposite direction. Instead of assuming that all notes belonging to the same denomination have the same visual appearance, the suggested approach treats old and new variants of a denomination as distinct classes during both dataset annotation and model training. In order to keep the entire system self-contained and appropriate for assistive use, this paper presents a currency denomination detection system based on YOLOv8n that follows an eleven-class scheme and explicitly separates old and new variants for applicable denominations. The detection system is combined with a lightweight offline text-to-speech engine that translates the detected denomination into spoken feedback without requiring an internet connection. The model is evaluated on a validation set of currency images, with special attention paid to per-class detection confidence, precision, recall, and inference latency on CPU-only hardware. The purpose of these measurements is to determine whether the suggested system can deliver accurate and fast enough performance for low-cost, real-time assistive deployment.

The remainder of this paper is organized as follows. Section II reviews related work on currency recognition for visually impaired users. Section III describes the dataset, model architecture, and training methodology. Section IV presents the experimental results. Finally, Section V concludes the paper by discussing the limitations of the proposed system and possible directions for future work.

II. RELATED WORKS

Many systems have been proposed to help blind or visually impaired users identify currency notes. Early methods relied on traditional image processing or simple CNN classifiers installed on smartphones or Raspberry Pi devices; more recent work has used modern object detectors.

YOLOv3-based systems achieved real-time banknote detection with reported accuracies around 95–96% on controlled test sets, usually paired with text-to-speech for audio feedback; YOLOv5 and YOLOv8 variants have further improved speed and accuracy on multi-denomination Indian and foreign currency datasets, frequently reaching mAP@0.5 values above 0.95 when trained on well-annotated collections.

A common limitation of these systems is that they concentrate on identifying a single note in isolation and stop at voice announcement of its denomination; relatively little attention has been paid to recognition under viewpoints typical of wearable, hands-free capture or to explicitly modeling the visual differences between old and new note designs for the same denomination, both of which are addressed in this work.

Some works combine currency recognition with obstacle detection or general object recognition on wearable platforms; offline text-to-speech libraries are commonly used to keep such systems operational. In addition to providing a useful wearable capture architecture that most previous recognition-only systems do not provide, this work expands on the robust recognition performance of lightweight YOLO models by explicitly differentiating between old and new note variants during dataset annotation and training.

III. METHODOLOGY

A. System Architecture

The system uses a capture-and-compute pipeline for hands-free operation, as shown in Fig. 1. A camera unit installed on smart glasses records a frame of the currency note and sends it to a companion smartphone application, which preprocesses the image (resizing and normalizing), passes the frame through a lightweight YOLO-based detector to identify the denomination label, and scores the prediction's confidence. The detected label is then mapped to a corresponding spoken phrase, synthesized using an on-device text-to-speech engine, and streamed to the user via Bluetooth-connected earphones or a speaker. Any reliance on network connectivity is eliminated by keeping the recognition and speech-synthesis phases fully on-device, which is crucial for a wearable assistive technology meant for constant, practical use.




Fig. 1. System architecture: a camera unit on smart glasses captures the note image, which is preprocessed and passed through the YOLO detector; the resulting label is converted to speech and streamed to the user over Bluetooth.

B. Data Collection and Preprocessing

Eleven classes were defined covering both new and old series notes: 10, 20 New, 20 Old, 50 New, 50 Old, 100 New, 100 Old, 200, 500 New, 2000, and the corresponding remaining variant. A representative training batch with bounding-box annotations is shown in Fig. 2, and the per-class instance counts along with the spatial distribution of the annotated boxes are summarized in Fig. 3.



Fig. 2. Representative augmented training batch with ground-truth bounding-box annotations.




Fig. 3. Dataset statistics: per-class instance counts (top-left), aggregate box positions (top-right), and normalized box center and size distributions (bottom row).

The dataset was augmented for real-world robustness (lighting variation, motion blur, partial

occlusion, reflective surfaces and crumpled notes). Identical preprocessing (resizing to 640×640 and normalization) is applied at both training and inference time so that the model sees consistent conditions in both stages.

C. Model Selection and Training

A YOLOv8n (nano) backbone was specifically chosen for deployability rather than peak accuracy alone. Training began with COCO-pretrained weights and ran for 100 epochs on CPU (batch size 16, image size 640). Standard Ultralytics augmentations (mosaic, HSV jitter, horizontal flip, scale, and translate) were used throughout training. Following training, the model was exported to the ONNX format (11.7 MB) for effective on-device inference. At inference time, a confidence threshold is applied, which causes a recapture request rather than informing the user of an uncertain result.

D. Evaluation Protocol

Detection metrics were calculated using a held-out validation set of 23 images; end-to-end latency is measured on the target hardware throughout the entire pipeline, including capture, preprocessing, inference, confidence check, speech synthesis, and audio playback; a separate real-world evaluation pass under live demo conditions (varying lighting, angles, and occlusion) is planned to report practical accuracy distinct from benchmark numbers.

IV. RESULTS AND DISCUSSION

A. Detection Performance

On the 23-image validation set the trained YOLOv8n model achieved:

- Precision: 0.98
- Recall: 1.00
- mAP@0.5: 0.995
- mAP@0.5:0.95: 0.91

All 11 classes reached mAP@0.5 = 0.995, as reflected in the precision-recall curve of Fig. 4. Per-class mAP@0.5:0.95 ranged from 0.846 (100 Old) to 0.995 (50 New).




Fig. 4. Precision-Recall curve for all eleven currency classes on the validation set.

Using the exported ONNX model, an independent detection test on the 23 validation images yielded a 100% detection rate (23/23), with minimum confidence of 0.932, maximum confidence of 0.990, and average confidence of 0.960. Fig. 5 displays representative qualitative detections, each annotated with the predicted class label and confidence score.



Fig. 5. Qualitative detection results on validation images, annotated with predicted denomination and confidence score.

B. Training Behaviour

There was no early stopping triggered during the entire 100 epochs of training. As seen in Fig. 6, box loss, classification loss, and DFL loss all steadily declined throughout training, but precision, recall, mAP@0.5, and mAP@0.5:0.95 all improved and stabilized in the later epochs. By epoch 11, mAP@0.5 reached the 0.995 ceiling and stayed steady, but mAP@0.5:0.95 kept improving to 0.91, suggesting that bounding-box precision was being refined rather than

overfitting. The risk of memorization is further decreased by the model's small capacity (roughly 3M parameters) in comparison to the 459 training images.




Fig. 6. Training and validation loss curves (box, classification, DFL) together with precision, recall, mAP@0.5 and mAP@0.5:0.95 over 100 training epochs.

Fig. 7 shows the confusion matrix on the validation set, normalized by true class. Confusion between different denominations and between old and new variations of the same denomination is minimal, as evidenced by the fact that all eleven classes are nearly entirely on the diagonal.




Fig. 7. Normalized confusion matrix on the validation set.

C. Inference Latency

Preprocessing took 1.7 ms, inference took 131 ms, and post-processing took 1.4 ms on the training CPU, for a total of roughly 134 ms per image. This number should significantly decrease on a smartphone with GPU or NPU acceleration, enabling real-time operation at several frames per second.

D. Voice Feedback

For each of the eleven denominations, the offline text-to-speech engine produced speech with success. The hands-free, low-latency design goal is achieved with audio generation latency of

about 0.5 seconds per phrase and no need for an internet connection.

E. Discussion

The findings show that a lightweight YOLOv8n model can achieve detection performance comparable to the literature while still being appropriate for mobile inference when trained on a carefully prepared dataset that explicitly distinguishes between old and new note variants. The high validation metrics and 100% detection rate on the test images attest to the successful recognition of Indian currency notes under the assessed conditions, including both their old and new design variants.

The current recognition pipeline does not address the open direction of extending the system to track cumulative transaction amounts, confirm that a physical note has left the capture frame, or otherwise confirm that a cash exchange has been fully completed.

The small validation set size, the absence of a fully held-out test split, and the still-limited evaluation under extreme real-world conditions (heavy occlusion, severe motion blur, worn notes) are limitations that need to be addressed in future work. Recognized gaps serve as the foundation for the upcoming experimental campaign.

F. Comparison with Existing Studies

In order to contextualize these findings, Table I contrasts the suggested system with representative currency-recognition studies published by IEEE, Springer, and other venues [1]–[8]. The comparison is indicative rather than a like-for-like benchmark because these studies differ in currency, dataset, model, and evaluation metric.

TABLE I
COMPARISON WITH EXISTING CURRENCY
RECOGNITION SYSTEMS

Ref.	Method	Currency data	and Reported result
[1]	YOLOv8 with gTTS	Counterfeit Indian notes	Fake-note check with voice output speech
[2]	YOLOv3	INR, 3,720 images, 7 classes	Detection 95.71%, recognition 100%
[3]	YOLOv8n with SE blocks	USD, EUR, BDT, 30 classes	mAP@0.5 97.21%, 3.1 ms per image
[4]	YOLOv9 on smart glasses	PKR, 3,611 images, 7 denominations	Accuracy up to 95%
[5]	CNN compared with YOLOv5	INR	CNN above 79.83% accuracy, about 1% below YOLOv5

Ref.	Method	Currency data	and Reported result
[6]	YOLOv5	INR, 10,000 images	Average accuracy 92.71%
[7]	IPCRNet (CNN classifier, 3.6M parameters)	INR, 50,263 images, 11 classes	Average accuracy 96.75%
[8]	Three-stage Faster CNN (VGG16)	KRW and JOD R- notes and coins	Multi-stage detection of notes and coins
This work	YOLOv8n (11.7 MB ONNX)	INR, 482 P 0.98, R 1.00, 11 mAP@0.5 0.995, 134 ms on CPU	

Table I makes three observations: First, the detection quality of the suggested model (precision 0.98, recall 1.00, mAP@0.5 of 0.995) is on par with or better than the results reported for the YOLO-based systems in [2], [3], and [6], which report a detection accuracy of 95.71%, an mAP@0.5 of 97.21%, and an average accuracy of 92.71%, respectively, and it also compares favorably with the 96.75% average accuracy of the lightweight classifier in [5] is also consistent with the use of a single-stage YOLO detector in place of the multi-stage pipeline of [8].

Second, the proposed system divides old and new note designs into eleven classes, while the YOLO-based Indian-currency systems [2] and [6] are based around seven denominations. Of the works compared here, [7] also models old and new notes as separate classes, but it is an image classifier and does not localize the note.

Third, the system is completely offline and lightweight. The YOLOv8n model has about 3M parameters and an 11.7 MB ONNX export, which is similar in size to the 3.6M-parameter network of [7] while also localizing the note. Speech is generated by an on-device engine, while [1] uses the gTTS library, which calls an online text-to-speech service. The end-to-end latency of about 134 ms per image reported in [3], but that difference may be due to hardware, and it is still sufficient for real-time use, as covered in Section IV-C.

The size of the current evaluation must be taken into consideration when reading these comparisons; for example, [7] is tested on approximately 9,500 images, while the suggested figures are obtained on a validation set of 23 images (Section IV-A). Therefore, a larger held-out test set is necessary before the aforementioned advantages can be deemed conclusive, as mentioned in Section IV-E.

V. CONCLUSION

In this paper, a real-time Indian currency denomination detection system for visually impaired users was presented. Precision 0.98, recall 1.00, mAP@0.5 0.995, and mAP@0.5:0.95 0.91 were attained by a YOLOv8n model trained on an 11-class dataset that clearly separates old and new note variants. The model was incorporated into a useful wearable design that includes a camera mounted on glasses and a companion smartphone that uses Bluetooth to provide offline text-to-speech feedback and inference.

The work demonstrates that precise, hands-free note recognition is possible on hardware with limited resources without compromising real-time performance. In order to improve the system's usefulness for separate financial transactions, future research will investigate expanding it to track cumulative transaction amounts and confirm full note exchange.

REFERENCES

- [1] R. Lalitha *et al.*, "Assisting visually challenged individuals to distinguish authentic currency from counterfeit," in *Advances in Data Science and Artificial Intelligence (ERCICA 2024)*, ser. Lecture Notes in Electrical Engineering, vol. 1399. Springer, 2025.
- [2] R. C. Joshi, S. Yadav, and M. K. Dutta, "YOLO-v3 based currency detection and recognition system for visually impaired persons," in *Proc. Int. Conf. Contemp. Comput. Appl. (IC3A)*, 2020, pp. 280–285.
- [3] S. A. Shreya, M. A. I. Siddique, and S. Tasnim, "Real-time currency detection and voice feedback for visually impaired individuals," arXiv:2510.20267, 2025.
- [4] M. Khalid, H. B. U. Haq, Z. Janjua, H. M. M. Akhtar, and H. M. Yousaf, "YOLOv9-based YOLO-enhanced smart glasses for real-time recognition of Pakistani currency: Empowering the visually impaired," *J. Comput. Biomed. Inform.*, vol. 8, no. 1, 2024.
- [5] S. D. Achar, C. S. Singh, C. S. S. Rao, K. P. Narayana, and A. Dasare, "Indian currency recognition system using CNN and comparison with YOLOv5," in *Proc. IEEE Int. Conf. Mobile Netw. Wireless Commun. (ICMNBC)*, 2021, pp. 1–6.
- [6] M. Pasumarthi, R. Padhy, R. Yadav, G. Subramaniam, and M. Rao, "An Indian currency recognition model for assisting visually impaired individuals," in *Proc. IEEE Int. Conf. Recent Adv. Syst. Sci. Eng. (RASSE)*, 2022, doi: 10.1109/RASSE54974.2022.9989624.
- [7] M. Singh, J. Chauhan, M. S. Kanroo, S. Verma, and P. Goyal, "IPCRF: An end-to-end Indian paper currency recognition framework for blind and visually impaired people," *IEEE Access*, vol. 10, pp. 90726–90744, 2022, doi: 10.1109/ACCESS.2022.3202007.
- [8] C. Park, S. W. Cho, N. R. Baek, J. Choi, and K. R. Park, "Deep feature-based three-stage detection of banknotes and coins for assisting visually impaired people," *IEEE Access*, vol. 8, pp. 184598–184613, 2020.

AI Detection Score



Page 2 of 7 - AI Writing Overview

Submission ID tnm:oid::3618:152472137

0% detected as AI

The percentage indicates the combined amount of likely AI-generated text as well as likely AI-generated text that was also likely revised by AI.

Caution: Review required.

It is essential to understand the limitations of AI detection before making decisions about a student's work. We encourage you to learn more about Turnitin's AI detection capabilities before using the tool.



0 AI-generated 0%

Likely generated or likely generated and revised by AI.

Disclaimer

Our AI writing assessment is designed to help educators identify text that might be prepared by a generative AI tool. Our AI writing assessment may not always be accurate (it may misidentify writing that is likely human generated as AI generated and likely AI generated as human generated) so it should not be used as the sole basis for adverse actions against a student. It takes further scrutiny and human judgment in conjunction with an organization's application of its specific academic policies to determine whether any academic misconduct has occurred.

Plagiarism Score



Page 2 of 9 - Integrity Overview

Submission ID trnoid::3618:152472137

14% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

Match Groups

-  **26 Not Cited or Quoted 11%**
Matches with neither in-text citation nor quotation marks
-  **3 Missing Quotations 2%**
Matches that are still very similar to source material
-  **2 Missing Citation 1%**
Matches that have quotation marks, but no in-text citation
-  **0 Cited and Quoted 0%**
Matches with in-text citation present, but no quotation marks

Top Sources

- 10%  Internet sources
- 9%  Publications
- 10%  Submitted works (Student Papers)

Integrity Flags

0 Integrity Flags for Review

No suspicious text manipulations found.

Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.

CONFERENCE/ JOURNAL PAPER PUBLICATION CURRENT SUBMISSION DETAILS

Conference Name: 2026 6th International Conference on Nanoelectronics, Machine Learning, Internet of Things & Computing Systems (NMIC)

STATUS: ACCEPTED

Track Name: NMIC2026 — Primary Subject Area: Artificial Intelligence

Paper ID: 2276

Created on: 2026-09-25 18:45:41 IST Last Modified: 2026-09-25 20:24:52 IST

Paper Title: An Edge-Deployable YOLOv8n-Based Indian Currency Recognition System with Offline Speech Assistance

Abstract: This paper presents a real-time currency denomination detection system based on the YOLOv8n object detector, designed to recognize not only different denominations but also old and new design variants of Indian currency notes, resulting in an eleven-class detection scheme. For visually impaired people, independently identifying the denomination of a currency note can be a persistent everyday challenge, especially when older and newer versions of the same denomination are in circulation because differences in size, color, and printed features can affect how reliably a recognition system can identify them. The model was trained on a custom-annotated dataset of 482 images and achieved a precision of 0.98, recall of 1.00, and [mAP@0.5](#) of 0.995 on the validation set. On CPU-only hardware, the system has a total inference latency of about 134 milliseconds per image, indicating its potential for low-cost deployment without dedicated GPU acceleration. An offline text-to-speech engine converts the detected denomination into spoken output, which is especially useful for assistive applications. Experimental testing on unseen images consistently produced high-confidence detections with an average of 0.96. These findings demonstrate that a lightweight model can achieve fine-grained recognition of currency note variants while maintaining real-time performance. As a result, the suggested system offers a workable and repeatable basis for assistive currency-recognition applications that can be implemented on edge devices with limited resources.

Authors:

1. CHETTIYAR VANI VIVEKANAND (chettiyarvanivivekanand@rajalakshmi.edu.in)
2. DHAANYA GK (230701069@rajalakshmi.edu.in)
3. FAREED AHAMED KM (230701087@rajalakshmi.edu.in)
4. SANJAY KRISHNA S (230701288@rajalakshmi.edu.in)
5. SAUNAK SEBASAN (230701298@rajalakshmi.edu.in)

Submission Files: phase1_paper_final_IEEE.pdf (536.9 Kb, 2026-09-25 18:40:43 IST)

PAPER ACCEPTANCE

Rajalakshmi Engineering College Mail - Acceptance Letter...

<https://mail.google.com/mail/u/1/?ik=082d847395&vi...>



SAUNAK RAMIYA SEBASAN 230701298 <230701298@raja...

Acceptance Letter of NMIC262276_CMT2276

2 messages

NMIC 2026 <nmic.2026.02@gmail.com>
To: 230701298@rajalakshmi.edu.in

Sun, Sep 27, 2026 at 12:28 PM

Acceptance Letter of NMIC262276_CMT2276

Dear Authors,

Congratulaons!

We are pleased to share that your arcle is accepted for presentaon and publicaon in Scopus IEEE Xplore NMIC-2026 with **minor revision**.

Download the registraon form from the given link and submit the filled form: https://isve.in/wpcontent/uploads/2025/11/Registration-Form_NMIC-2025.pdf within 3-5 days of this mail. Arcle must be prepared in IEEE double column format: <https://isve.in/wp-content/uploads/2025/11/IEEEconference-template-a4.docx> . Aer incorporaon of the reviewers comments, submit your final arcle DOCX, PDF and Plagiarism report in the name of your paper ID (e.g. NMIC26001_CMT1) to reply to this mail nmic.2026.02@gmail.com where the PDF file must be passed through IEEE PDF Express available at CMT Portal and also upload on CMT portal against your submission ID: [https:// cmt3.research.microsoft.com/NMIC2026](https://cmt3.research.microsoft.com/NMIC2026). Registraon has been started whose arcle reviews are completed and communicated to authors. If facing any problem to convert PDF from IEEE PDF express submit docx file.

Documents required for Registraon: Registraon Form, Proof of Payments, Arcle DOCX, PDF, Plagiarism Report etc. Copyright Form will be submied later. It will be informed by the commiee in the due course of me.

Registraon informaon details: <https://isve.in/ieee-nmic-2026-registration-details/>

Conference informaon details: <https://isve.in/nmic-2026/>

Benefits of IEEE/ ISVE Membership: Authors will get 5% discount in their arcle fees aer the Membership of IEEE/ ISVE. ISVE Membership

details are available: <https://isve.in/isvemembership/> interested can fill and submit the form as per their eligibility.

Reviewer comments are given below:-

1. The idea of the article Secure-GPT: Detecting Sensitive Data Leakage in Browser-Based LLM Chat Interfaces is fine.
2. Author affiliations are given.
3. Number of references is fine.
4. Comparative studies with standard publications like IEEE/ Springer etc. are not shown. Highlight your findings.
5. Formatting of the article is fine.
6. The Figures & Tables numbering and captions are given.

1 of 2 04-10-2026, 06:23 pm Rajalakshmi Engineering College Mail - Acceptance Letter...

<https://mail.google.com/mail/u/1/?ik=082d847395&vi...>

7. Simulation results are shown and explained.
8. References are properly formatted and aligned as per IEEE.
9. Please note that plagiarism should not be more than 10%. Please submit the final paper with a plagiarism report.
10. Send the updated docx and PDF revised file as camera ready paper in the given format
<https://isve.in/wpcontent/uploads/2025/11/IEEE-conference-template-a4.docx>
11. Articles will be accepted with the above changes.

Thanks & regards
